

## “Capítol 24: Advanced Application Development”

Fitxers i bases de dades

## Chapter 12: Advanced Application Development

- Performance Tuning
- Performance Benchmarks
- Standardization
- E-Commerce
- Legacy Systems

# Performance Tuning

- Adjusting various parameters and design choices to improve system performance for a specific application.
- Tuning is best done by
  - 1 identifying bottlenecks, and
  - 2 eliminating them.
- Can tune a database system at 3 levels:
  - **Hardware** - e.g., add disks to speed up I/O, add memory to increase buffer hits, move to a faster processor.
  - **Database system parameters** - e.g., set buffer size to avoid paging of buffer, set checkpointing intervals to limit log size. System may have automatic tuning.
  - **Higher level database design**, such as the schema, indices and transactions (more later)

# Bottlenecks

- Performance of most systems (at least before they are tuned) usually limited by performance of one or a few components: these are called **bottlenecks**
  - E.g., 80% of the code may take up 20% of time and 20% of code takes up 80% of time  
Worth spending most time on 20% of code that take 80% of time
- Bottlenecks may be in hardware (e.g., disks are very busy, CPU is idle), or in software
- Removing one bottleneck often exposes another
- De-bottlenecking consists of repeatedly finding bottlenecks, and removing them
  - This is a heuristic

# Identifying Bottlenecks

- Transactions request a sequence of services
  - E.g., CPU, Disk I/O, locks
- With concurrent transactions, transactions may have to wait for a requested service while other transactions are being served
- Can model database as a **queueing system** with a queue for each service
  - Transactions repeatedly do the following
    - ▶ request a service, wait in queue for the service, and get serviced
- Bottlenecks in a database system typically show up as very high utilizations (and correspondingly, very long queues) of a particular service
  - E.g., disk vs. CPU utilization
  - 100% utilization leads to very long waiting time:
    - ▶ Rule of thumb: design system for about 70% utilization at peak load
    - ▶ utilization over 90% should be avoided

# Queues In A Database System

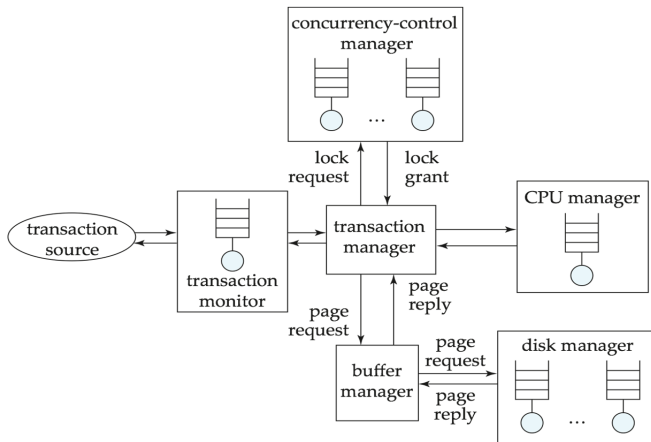


Figure: Example.

# Tunable Parameters

- Tuning of hardware
- Tuning of schema
- Tuning of indices
- Tuning of materialized views
- Tuning of transactions

# Tuning of Hardware

- Even well-tuned transactions typically require a few I/O operations
  - Typical disk supports about 100 random I/O operations per second
  - Suppose each transaction requires just 2 random I/O operations. Then to support  $n$  transactions per second, we need to stripe data across  $n/50$  disks (ignoring skew)
- Number of I/O operations per transaction can be reduced by keeping more data in memory
  - If all data is in memory, I/O needed only for writes
  - Keeping frequently used data in memory reduces disk accesses, reducing number of disks required, but has a memory cost



# Hardware Tuning: Five-Minute Rule

- Question: which data to keep in memory:
  - If a page is accessed  $n$  times per second, keeping it in memory saves
    - ▶  $n * \frac{\text{price-per-disk-drive}}{\text{accesses-per-second-per-disk}}$
  - Cost of keeping page in memory
    - ▶  $\frac{\text{price-per-MB-of-memory}}{\text{ages-per-MB-of-memory}}$
  - Break-even point: value of  $n$  for which above costs are equal
    - ▶ If accesses are more then saving is greater than cost
  - Solving above equation with current disk and memory prices leads to:  
**5-minute rule: if a page that is randomly accessed is used more frequently than once in 5 minutes it should be kept in memory**
    - ▶ (by buying sufficient memory!)

# Hardware Tuning: One-Minute Rule

- For sequentially accessed data, more pages can be read per second.  
Assuming sequential reads of 1MB of data at a time: **1-minute rule: sequentially accessed data that is accessed once or more in a minute should be kept in memory**
- Prices of disk and memory have changed greatly over the years, but the ratios have not changed much
  - So rules remain as 5 minute and 1 minute rules, not 1 hour or 1 second rules!

# Hardware Tuning: Choice of RAID Level

- To use RAID 1 or RAID 5?
  - Depends on ratio of reads and writes
    - ▶ RAID 5 requires 2 block reads and 2 block writes to write out one data block
- If an application requires  $r$  reads and  $w$  writes per second
  - RAID 1 requires  $r + 2w$  I/O operations per second
  - RAID 5 requires:  $r + 4w$  I/O operations per second
- For reasonably large  $r$  and  $w$ , this requires lots of disks to handle workload
  - RAID 5 may require more disks than RAID 1 to handle load!
  - Apparent saving of number of disks by RAID 5 (by using parity, as opposed to the mirroring done by RAID 1) may be illusory!
- Thumb rule: RAID 5 is fine when writes are rare and data is very large, but RAID 1 is preferable otherwise
  - If you need more disks to handle I/O load, just mirror them since disk capacities these days are enormous!

# Tuning the Database Design

- Schema tuning

- Vertically partition relations to isolate the data that is accessed most often
  - only fetch needed information.
    - ▶ E.g., split account into two, (account-number, branch-name) and (account-number, balance).
      - Branch-name need not be fetched unless required
- Improve performance by storing a **denormalized relation**
  - E.g., store join of account and depositor; branch-name and balance information is repeated for each holder of an account, but join need not be computed repeatedly.
    - Price paid: more space and more work for programmer to keep relation consistent on updates
  - Better to use materialized views (more on this later..)
- Cluster together on the same disk page records that would match in a frequently required join
  - ▶ Compute join very efficiently when required.

# Tuning the Database Design (Cont.)

- Index tuning
  - Create appropriate indices to speed up slow queries/updates
  - Speed up slow updates by removing excess indices (tradeoff between queries and updates)
  - Choose type of index (B-tree/hash) appropriate for most frequent types of queries.
  - Choose which index to make clustered
- Index tuning wizards look at past history of queries and updates ([the workload](#)) and recommend which indices would be best for the workload

# Tuning the Database Design (Cont.)

## Materialized Views

- Materialized views can help speed up certain queries
  - Particularly aggregate queries
- Overheads
  - Space
  - Time for view maintenance
    - ▶ Immediate view maintenance: done as part of update txn
      - time overhead paid by update transaction
    - ▶ Deferred view maintenance: done only when required
      - Update transaction is not affected, but system time is spent on view maintenance
        - Until updated, the view may be out-of-date
- Preferable to denormalized schema since view maintenance is systems responsibility, not programmers
  - Avoids inconsistencies caused by errors in update programs

# Tuning the Database Design (Cont.)

- How to choose set of materialized views
  - Helping one transaction type by introducing a materialized view may hurt others
  - Choice of materialized views depends on costs
    - ▶ Users often have no idea of actual cost of operations
  - Overall, manual selection of materialized views is tedious
- Some database systems provide tools to help DBA choose views to materialize
  - “Materialized view selection wizards”

# Tuning of Transactions

- Basic approaches to tuning of transactions
  - Improve set orientation
  - Reduce lock contention
- Rewriting of queries to improve performance was important in the past, but smart optimizers have made this less important
- Communication overhead and query handling overheads significant part of cost of each call
  - Combine multiple embedded SQL/ODBC/JDBC queries into a single set-oriented query
    - ▶ Set orientation → fewer calls to database
    - ▶ E.g., tune program that computes total salary for each department using a separate SQL query by instead using a single query that computes total salaries for all department at once (using **group by**)
  - Use stored procedures: avoids re-parsing and re-optimization of query



## Tuning of Transactions (Cont.)

- Reducing lock contention
- Long transactions (typically read-only) that examine large parts of a relation result in lock contention with update transactions
  - E.g., large query to compute bank statistics and regular bank transactions
- To reduce contention
  - Use multi-version concurrency control
    - ▶ E.g., Oracle “snapshots” which support multi-version 2PL
  - Use degree-two consistency (cursor-stability) for long transactions
    - ▶ Drawback: result may be approximate

## Tuning of Transactions (Cont.)

- Long update transactions cause several problem
  - Exhaust lock space
  - Exhaust log space
    - ▶ and also greatly increase recovery time after a crash, and may even exhaust log space during recovery if recovery algorithm is badly designed!
- Use **mini-batch** transactions to limit number of updates that a single transaction can carry out. E.g., if a single large transaction updates every record of a very large relation, log may grow too big.
  - Split large transaction into batch of "mini-transactions", each performing part of the updates
    - Hold locks across transactions in a mini-batch to ensure serializability
      - ▶ If lock table size is a problem can release locks, but at the cost of serializability
  - In case of failure during a mini-batch, must complete its remaining portion on recovery, to ensure atomicity.

# Performance Simulation

- **Performance simulation** using queuing model useful to predict bottlenecks as well as the effects of tuning changes, even without access to real system
- Queuing model as we saw earlier
  - Models activities that go on in parallel
- Simulation model is quite detailed, but usually omits some low level details
  - Model **service time**, but disregard details of service
  - E.g., approximate disk read time by using an average disk read time
- Experiments can be run on model, and provide an estimate of measures such as average throughput/response time
- Parameters can be tuned in model and then replicated in real system
  - E.g., number of disks, memory, algorithms, etc.

# Performance Benchmarks

- Suites of tasks used to quantify the performance of software systems
- Important in comparing database systems, especially as systems become more standards compliant.
- Commonly used performance measures:
  - **Throughput** (transactions per second, or tps)
  - **Response time** (delay from submission of transaction to return of result)
  - **Availability** or mean time to failure

## Performance Benchmarks (Cont.)

- Suites of tasks used to characterize performance
  - single task not enough for complex systems
- Beware when computing average throughput of different transaction types
  - E.g., suppose a system runs transaction type A at 99 tps and transaction type B at 1 tps.
  - Given an equal mixture of types A and B, throughput is not  $(99+1)/2 = 50$  tps.
  - Running one transaction of each type takes time  $1+.01$  seconds, giving a throughput of 1.98 tps.
  - To compute average throughput, use **harmonic mean**:

$$\frac{n}{\frac{1}{t_1} + \frac{1}{t_2} + \dots + \frac{1}{t_n}}$$

- **Interference** (e.g., lock contention) makes even this incorrect if different transaction types run concurrently

# Database Application Classes

- Online transaction processing (OLTP)
  - requires high concurrency and clever techniques to speed up commit processing, to support a high rate of update transactions.
- Decision support applications
  - including online analytical processing, or OLAP applications
  - require good query evaluation algorithms and query optimization.
- Architecture of some database systems tuned to one of the two classes
  - E.g., Teradata is tuned to decision support
- Others try to balance the two requirements
  - E.g., Oracle, with snapshot support for long read-only transaction

- The Transaction Processing Council (TPC) benchmark suites are widely used.
  - **TPC-A** and **TPC-B**: simple OLTP application modeling a bank teller application with and without communication
    - ▶ Not used anymore
  - **TPC-C**: complex OLTP application modeling an inventory system
    - ▶ Current standard for OLTP benchmarking

## Benchmarks Suites (Cont.)

- TPC benchmarks (cont.)
  - **TPC-D:** complex decision support application
    - ▶ Superceded by TPC-H and TPC-R
  - **TPC-H:** (H for ad hoc) based on TPC-D with some extra queries
    - ▶ Models ad hoc queries which are not known beforehand
      - Total of 22 queries with emphasis on aggregation
    - ▶ prohibits materialized views
    - ▶ permits indices only on primary and foreign keys
  - **TPC-R:** (R for reporting) same as TPC-H, but without any restrictions on materialized views and indices
  - **TPC-W:** (W for Web) End-to-end Web service benchmark modeling a Web bookstore, with combination of static and dynamically generated pages



# TPC Performance Measures

- TPC performance measures
  - **transactions-per-second** with specified constraints on response time
  - **transactions-per-second-per-dollar** accounts for cost of owning system
- TPC benchmark requires database sizes to be scaled up with increasing transactions-per-second
  - Reflects real world applications where more customers means more database size and more transactions-per-second
- External audit of TPC performance numbers mandatory
  - TPC performance claims can be trusted

# TPC Performance Measures

- Two types of tests for TPC-H and TPC-R

- **Power test**: runs queries and updates sequentially, then takes mean to find queries per hour
- **Throughput test**: runs queries and updates concurrently
  - ▶ multiple streams running in parallel each generates queries, with one parallel update stream
- **Composite query per hour metric**: square root of product of power and throughput metrics
- **Composite price/performance metric**

## Other Benchmarks

- OODB transactions require a different set of benchmarks.
  - OO7 benchmark has several different operations, and provides a separate benchmark number for each kind of operation
  - Reason: hard to define what is a typical OODB application
- Benchmarks for XML being discussed

# Standardization

- The complexity of contemporary database systems and the need for their interoperability require a variety of standards.
  - syntax and semantics of programming languages
  - functions in application program interfaces
  - data models (e.g., object oriented/object relational databases)
- **Formal standards** are standards developed by a standards organization (ANSI, ISO), or by industry groups, through a public process
- **De facto standards** are generally accepted as standards without any formal process of recognition
  - Standards defined by dominant vendors (IBM, Microsoft) often become de facto standards
  - De facto standards often go through a formal process of recognition and become formal standards

## Standardization (Cont.)

- **Anticipatory standards** lead the market place, defining features that vendors then implement .
  - Ensure compatibility of future products
  - But at times become very large and unwieldy since standards bodies may not pay enough attention to ease of implementation (e.g., SQL-92 or SQL:1999)
- **Reactionary standards** attempt to standardize features that vendors have already implemented, possibly in different ways.
  - Can be hard to convince vendors to change already implemented features. E.g., OODB systems

# SQL Standards History

- SQL developed by IBM in late 70s/early 80s
- SQL-86 first formal standard
- IBM SAA standard for SQL in 1987
- SQL-89 added features to SQL-86 that were already implemented in many systems
  - Was a reactionary standard
- SQL-92 added many new features to SQL-89 (anticipatory standard)
  - Defines levels of compliance (entry, intermediate and full)
  - Even now few database vendors have full SQL-92 implementation

# SQL Standards History (Cont.)

- SQL:1999

- Adds variety of new features - extended data types, object orientation, procedures, triggers, etc.
- Broken into several parts
  - ▶ SQL/Framework (Part 1): overview
  - ▶ SQL/Foundation (Part 2): types, schemas, tables, query/update statements, security, etc.
  - ▶ SQL/CLI (Call Level Interface) (Part 3): API interface
  - ▶ SQL/PSM (Persistent Stored Modules) (Part 4): procedural extensions
  - ▶ SQL/Bindings (Part 5): embedded SQL for different embedding languages

## SQL Standards History (Cont.)

- More parts undergoing standardization process
  - Part 7: SQL/Temporal: temporal data
  - Part 9: SQL/MED (Management of External Data)
    - ▶ Interfacing of database to external data sources
      - Allows other databases, even files, can be viewed as part of the database
  - Part 10 SQL/OLB (Object Language Bindings): embedding SQL in Java
  - Missing part numbers 6 and 8 cover features that are not near standardization yet



## Database Connectivity Standards.)

- **Open DataBase Connectivity (ODBC)**: standard for database interconnectivity
  - based on Call Level Interface (CLI) developed by X/Open consortium
  - defines application programming interface, and SQL features that must be supported at different levels of compliance
- **JDBC** standard used for Java
- **X/Open XA** standards define transaction management standards for supporting distributed 2-phase commit
- **OLE-DB**: API like ODBC, but intended to support non-database sources of data such as flat files
  - OLE-DB program can negotiate with data source to find what features are supported
  - Interface language may be a subset of SQL
- **ADO (Active Data Objects)**: easy-to-use interface to OLE-DB functionality

# Object Oriented Databases Standards.

- **Object Database Management Group (ODMG)** standard for object-oriented databases
  - version 1 in 1993 and version 2 in 1997, version 3 in 2000
  - provides language independent Object Definition Language (ODL) as well as several language specific bindings
- **Object Management Group (OMG)** standard for distributed software based on objects
  - **Object Request Broker (ORB)** provides transparent message dispatch to distributed objects
  - **Interface Definition Language (IDL)** for defining language-independent data types
  - **Common Object Request Broker Architecture (CORBA)** defines specifications of ORB and IDL

## XML-Based Standards.

- Several XML based Standards for E-commerce
  - E.g., RosettaNet (supply chain), BizTalk
  - Define catalogs, service descriptions, invoices, purchase orders, etc.
  - XML wrappers are used to export information from relational databases to XML
- Simple Object Access Protocol (SOAP): XML based remote procedure call standard
  - Uses XML to encode data, HTTP as transport protocol
  - Standards based on SOAP for specific applications
    - ▶ E.g., OLAP and Data Mining standards from Microsoft

# E-Commerce.

- E-commerce is the process of carrying out various activities related to commerce through electronic means
- Activities include:
  - Presale activities: catalogs, advertisements, etc.
  - Sale process: negotiations on price/quality of service
  - Marketplace: e.g., stock exchange, auctions, reverse auctions
  - Payment for sale
  - Delivery related activities: electronic shipping, or electronic tracking of order processing/shipping
  - Customer support and post-sale service

# E-Catalogs.

- Product catalogs must provide searching and browsing facilities
  - Organize products into intuitive hierarchy
  - Keyword search
  - Help customer with comparison of products
- Customization of catalog:
  - Negotiated pricing for specific organizations
  - Special discounts for customers based on past history
    - ▶ E.g., loyalty discount
  - Legal restrictions on sales
    - ▶ Certain items not exposed to under-age customers
- Customization requires extensive customer-specific information

# Marketplaces.

- Marketplaces help in negotiating the price of a product when there are multiple sellers and buyers
- Several types of marketplaces
  - Reverse auction
  - Auction
  - Exchange
- Real world marketplaces can be quite complicated due to product differentiation
- Database issues:
  - Authenticate bidders
  - Record buy/sell bids securely
  - Communicate bids quickly to participants
    - ▶ Delays can lead to financial loss to some participants
  - Need to handle very large volumes of trade at times
    - ▶ E.g., at the end of an auction

# Types of Marketplace.

- **Reverse auction system:** single buyer, multiple sellers.
  - Buyer states requirements, sellers bid for supplying items. Lowest bidder wins. (also known as tender system)
  - **Open bidding** vs. **closed bidding**
- **Auction:** Multiple buyers, single seller
  - Simplest case: only one instance of each item is being sold
  - Highest bidder for an item wins
  - More complicated with multiple copies, and buyers bid for specific number of copies
- **Exchange:** multiple buyers, multiple sellers
  - E.g., stock exchange
  - Buyers specify maximum price, sellers specify minimum price
  - exchange matches buy and sell bids, deciding on price for the trade
    - ▶ e.g., average of buy/sell bids

# Order Settlement.

- Order settlement: payment for goods and delivery
- Insecure means for electronic payment: send credit card number
  - Buyers may present some one else's credit card numbers
  - Seller has to be trusted to bill only for agreed-on item
  - Seller has to be trusted not to pass on the credit card number to unauthorized people
- Need secure payment systems
  - Avoid above-mentioned problems
  - Provide greater degree of privacy
    - ▶ E.g., not reveal buyers identity to seller
  - Ensure that anyone monitoring the electronic transmissions cannot access critical information



# Secure Payment Systems.

- All information must be encrypted to prevent eavesdropping
  - Public/private key encryption widely used
- Must prevent **person-in-the-middle attacks**
  - E.g., someone impersonates seller or bank/credit card company and fools buyer into revealing information
    - ▶ Encrypting messages alone doesn't solve this problem
    - ▶ More on this in next slide
- Three-way communication between seller, buyer and credit-card company to make payment
  - Credit card company credits amount to seller
  - Credit card company consolidates all payments from a buyer and collects them together
    - ▶ E.g., via buyer's bank through physical/electronic check payment

## Secure Payment Systems (Cont).

- **Digital certificates** are used to prevent impersonation/man-in-the middle attack
  - Certification agency creates digital certificate by encrypting, e.g., seller's public key using its own private key
    - ▶ Verifies sellers identity by external means first!
  - Seller sends certificate to buyer
  - Customer uses public key of certification agency to decrypt certificate and find sellers public key
    - ▶ Man-in-the-middle cannot send fake public key
  - Sellers public key used for setting up secure communication
- Several secure payment protocols
  - E.g., Secure Electronic Transaction (SET)

# Digital Cash.

- Credit-card payment does not provide anonymity
  - The SET protocol hides buyers identity from seller
  - But even with SET, buyer can be traced with help of credit card company
- Digital cash systems provide anonymity similar to that provided by physical cash
  - E.g., [Dig Cash](#)
  - Based on encryption techniques that make it impossible to find out who purchased digital cash from the bank
  - Digital cash can be spent by purchaser in parts
    - ▶ much like writing a check on an account whose owner is anonymous

# Legacy Systems.

- Legacy systems are older-generation systems that are incompatible with current generation standards and systems but still in production use
  - E.g., applications written in Cobol that run on mainframes
    - ▶ Today's hot new system is tomorrows legacy system!
- Porting legacy system applications to a more modern environment is problematic
  - Very expensive, since legacy system may involve millions of lines of code, written over decades
    - ▶ Original programmers usually no longer available
  - Switching over from old system to new system is a problem
    - ▶ more on this later
- One approach: build a **wrapper** layer on top of legacy application to allow interoperation between newer systems and legacy application
  - E.g., use ODBC or OLE-DB as wrapper

## Legacy Systems (Cont.)

- Rewriting legacy application requires a first phase of understanding what it does
  - Often legacy code has no documentation or outdated documentation
  - **reverse engineering**: process of going over legacy code to
    - ▶ Come up with schema designs in ER or OO model
    - ▶ Find out what procedures and processes are implemented, to get a high level view of system
- **Re-engineering**: reverse engineering followed by design of new system
  - Improvements are made on existing system design in this process

## Legacy Systems (Cont.)

- Switching over from old to new system is a major problem
  - Production systems are in every day, generating new data
  - Stopping the system may bring all of a company's activities to a halt, causing enormous losses
- Big-bang approach:
  - 1 Implement complete new system
  - 2 Populate it with data from old system
    - 1 No transactions while this step is executed
    - 2 scripts are created to do this quickly
  - 3 Shut down old system and start using new system
  - **Danger with this approach:** what if new code has bugs or performance problems, or missing features
    - Company may be brought to a halt

## Legacy Systems (Cont.)

- Chicken-little approach:

- Replace legacy system one piece at a time
- Use wrappers to interoperate between legacy and new code
  - ▶ E.g., replace front end first, with wrappers on legacy backend
    - Old front end can continue working in this phase in case of problems with new front end
  - ▶ Replace back end, one functional unit at a time
    - All parts that share a database may have to be replaced together, or wrapper is needed on database also
- Drawback: significant extra development effort to build wrappers and ensure smooth interoperation
  - ▶ Still worth it if company's life depends on system

## FINAL DEL CAPÍTULO 24